

# Automatic Differentiation

COMP6211J

Binhang Yuan

# Numerical Differentiation

# Numerical Differentiation

- Numerical differentiation is the finite difference approximation of derivatives using values of the original function evaluated at some sample points.

- It is based on the limit definition of a derivative of function

$f: \mathbb{R}^n \rightarrow \mathbb{R} :$


$$\frac{\partial f}{\partial x_i} = \lim_{\epsilon \rightarrow 0} \frac{f(\mathbf{x} + \epsilon \mathbf{e}_i) - f(\mathbf{x})}{\epsilon} \approx \frac{f(\mathbf{x} + h \mathbf{e}_i) - f(\mathbf{x})}{h}$$

- $\mathbf{e}_i$  is the  $i$ -th unit vector,  $h > 0$  is a small step size.

# Pros and Cons

- Advantage:
  - Easy to implement.
- Disadvantage:
  - Perform  $\mathcal{O}(n)$  evaluations of  $f$  for a gradient in  $n$  dimensions.
  - Requires careful consideration in selecting the step size  $h$ .

# Choose Step Size $h$



- Truncation Error:
  - The error of approximation that one gets from  $h$  not actually being zero.
  - Proportional to a power of  $h$ .
- Rounding Error:
  - The inaccuracy that is inflicted by the limited precision of computations.
  - Inversely proportional to a power of  $h$ .

# Symbolic Differentiation

# Derivative Computation Rules

- Assume  $f(x): \mathbb{R} \rightarrow \mathbb{R}$ ,  $g(x): \mathbb{R} \rightarrow \mathbb{R}$ :
- Derivative of sum or difference:  $u = f(x), v = g(x)$ :

$$\frac{d}{dx}(u \pm v) = \frac{du}{dx} \pm \frac{dv}{dx}$$

- Product Rule:  $u = f(x), v = g(x)$ :

$$\frac{d}{dx}(uv) = u \frac{dv}{dx} + v \frac{du}{dx}$$

- Chain Rule:  $y = f(u), u = g(x)$ :

$$\frac{dy}{dx} = \frac{dy}{du} \frac{du}{dx}$$

# Derivative of Common Functions

- $f(x) = c, \quad \frac{df(x)}{dx} = 0$
- $f(x) = x, \quad \frac{df(x)}{dx} = 1$
- $f(x) = cx, \quad \frac{df(x)}{dx} = c$
- $f(x) = x^n, \quad \frac{df(x)}{dx} = nx^{n-1}$
- $f(x) = e^x, \quad \frac{df(x)}{dx} = e^x$

- $f(x) = \ln(x), \quad \frac{df(x)}{dx} = \frac{1}{x}$
- $f(x) = \sin(x), \quad \frac{df(x)}{dx} = \cos(x)$
- $f(x) = \cos(x), \quad \frac{df(x)}{dx} = -\sin(x)$
- $f(x) = \tan(x), \quad \frac{df(x)}{dx} = \sec^2(x)$



# Main Idea

- Symbolic differentiation is the automatic manipulation of expressions for obtaining derivative expressions carried out by applying derivative computation rules.
- When formulae are represented as data structures, symbolically differentiating an expression tree is a perfectly mechanistic process.
- This is realized in modern computer algebra systems such as Mathematica.

# Problem

- Symbolic derivatives do not lend themselves to efficient runtime calculation of derivative values, as they can get exponentially larger than the expression whose derivative they represent.
- Expression swell: careless symbolic differentiation can easily produce exponentially large symbolic expressions that take correspondingly long to evaluate.

# Expression Swell

Iterations of the logistic map  $l_{n+1} = 4l_n(1 - l_n)$ ,  $l_1 = x$  and the corresponding derivatives of  $l_n$  with respect to  $x$ , illustrating expression swell.

$n$	$l_n$	$\frac{d}{dx}l_n$	$\frac{d}{dx}l_n$ (Simplified form)
1	$x$	1	1
2	$4x(1 - x)$	$4(1 - x) - 4x$	$4 - 8x$
3	$16x(1 - x)(1 - 2x)^2$	$16(1 - x)(1 - 2x)^2 - 16x(1 - 2x)^2 - 64x(1 - x)(1 - 2x)$	$16(1 - 10x + 24x^2 - 16x^3)$
4	$64x(1 - x)(1 - 2x)^2(1 - 8x + 8x^2)^2$	$128x(1 - x)(-8 + 16x)(1 - 2x)^2(1 - 8x + 8x^2) + 64(1 - x)(1 - 2x)^2(1 - 8x + 8x^2)^2 - 64x(1 - 2x)^2(1 - 8x + 8x^2)^2 - 256x(1 - x)(1 - 2x)(1 - 8x + 8x^2)^2$	$64(1 - 42x + 504x^2 - 2640x^3 + 7040x^4 - 9984x^5 + 7168x^6 - 2048x^7)$

# Automatic Differentiation

# Main Idea

- An automatic differentiation (AD) system will convert the program into a sequence of elementary operations with specified routines for computing derivatives:
  - Apply symbolic differentiation at the elementary operation level;
  - Keep intermediate numerical results;
  - Combining the derivatives of the constituent operations through the chain rule gives the derivative of the overall composition.

# Notations

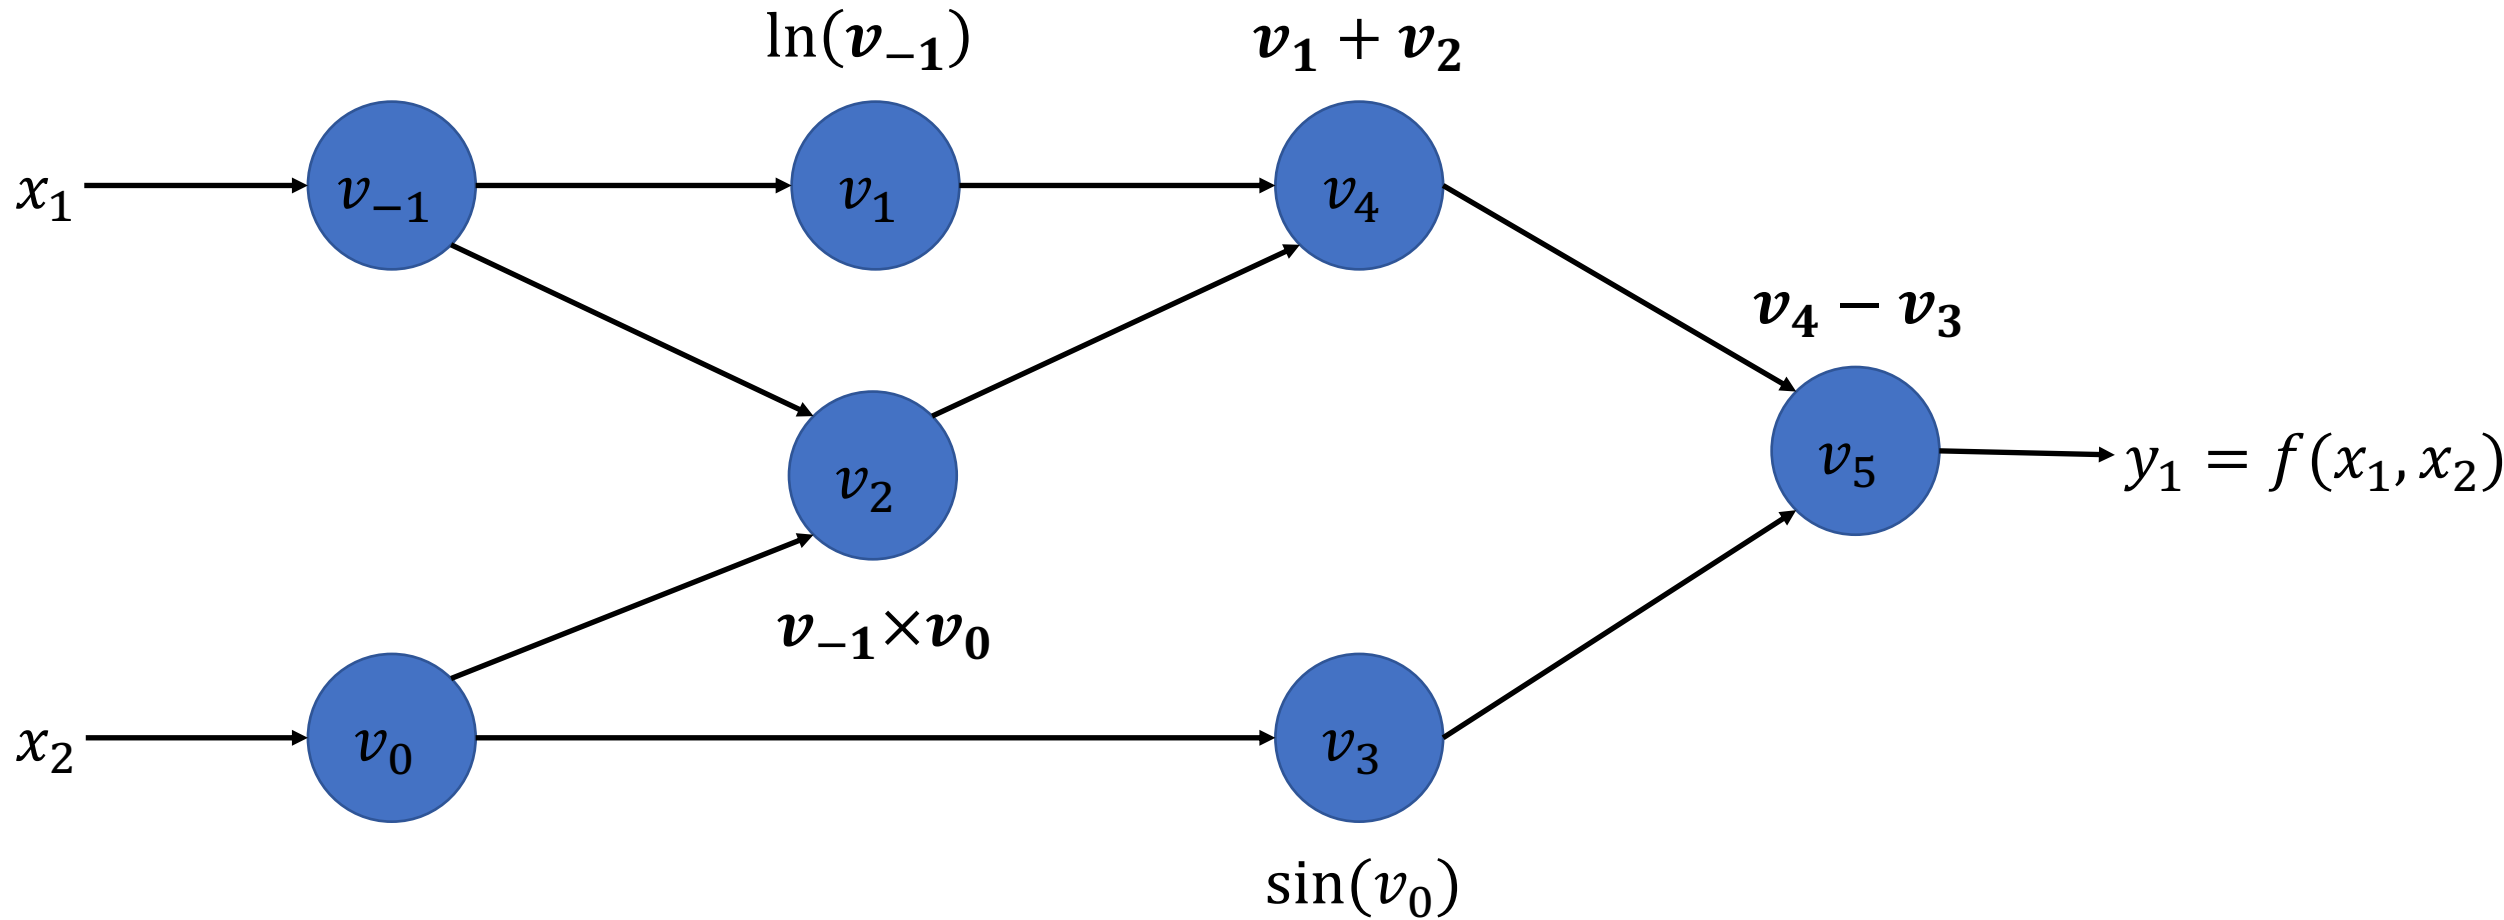
- The Jacobian matrix of a function  $f: \mathbb{R}^n \rightarrow \mathbb{R}^m$  is defined by a  $m \times n$  matrix noted by  $\mathbf{J}$  where  $J_{ij} = \frac{\partial y_i}{\partial x_j}$ , or explicitly:

$$\mathbf{J} = \begin{bmatrix} \frac{\partial y_1}{\partial x_1} & \dots & \frac{\partial y_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial y_m}{\partial x_1} & \dots & \frac{\partial y_m}{\partial x_n} \end{bmatrix}$$

# Notations

- A function  $f: \mathbb{R}^n \rightarrow \mathbb{R}^m$  is constructed using intermediate variable  $v_i$  such that:
  - Variable  $v_{j-n} = x_j$ ,  $j = 1, \dots, n$ , are the input variables;
  - Variable  $v_i$ ,  $i = 1, \dots, l$ , are the intermediate variables;
  - Variable  $y_{m-k} = v_{l-k}$ ,  $k = 1, \dots, m$ , are the output variables;

Example:  $f(x_1, x_2) = \ln(x_1) + x_1x_2 - \sin(x_2)$





# Forward Mode AD


- For computing the derivative of  $f$  with respect to  $x_1$ , we start by associating with each intermediate variable  $v_i$  a derivative (tangent):

$$\dot{v}_i = \frac{\partial v_i}{\partial x_1}$$

- Apply the chain rule to each elementary operation in the forward primal trace;
- Generate the corresponding tangent (derivative) trace;
- Evaluating the primals  $v_i$  in lockstep with their corresponding tangents  $\dot{v}_i$  gives us the required derivative in the final variable  $\dot{v}_5 = \frac{\partial y_1}{\partial x_1}$ .

# Forward Mode AD:

$$f(x_1, x_2) = \ln(x_1) + x_1 x_2 - \sin(x_2)$$



Forward Primal Trace

$v_{-1} = x_1$	$= 2$
$v_0 = x_2$	$= 5$
$v_1 = \ln(v_{-1})$	$= \ln(2) = 0.693$
$v_2 = v_{-1} \times v_0$	$= 2 \times 5 = 10$
$v_3 = \sin v_0$	$= \sin 5 = 0.959$
$v_4 = v_1 + v_2$	$= 0.693 + 10$
$v_5 = v_4 - v_3$	$= 10.693 + 0.959$
$y_1 = v_5$	$= 11.652$

Forward Tangent (Derivative) Trace

$\dot{v}_{-1} = \dot{x}_1$	$= 1$	$\dot{v}_{-1} = \frac{\partial x_1}{\partial x_1} = 1$
$\dot{v}_0 = \dot{x}_2$	$= 0$	
$\dot{v}_1 = \dot{v}_{-1} / v_{-1}$	$= 1/2$	
$\dot{v}_2 = \dot{v}_{-1} \times v_0 + \dot{v}_0 \times v_{-1}$	$= 1 \times 5 + 0 \times 2$	
$\dot{v}_3 = \dot{v}_0 \times \cos v_0$	$= 0 \times \cos 5$	
$\dot{v}_4 = \dot{v}_1 + \dot{v}_2$	$= 0.5 + 5$	
$\dot{v}_5 = \dot{v}_4 - \dot{v}_3$	$= 5.5 - 0$	
$\dot{y}_1 = \dot{v}_5$	$= 5.5$	

# Forward Mode AD

- Compute the Jacobian of a function  $f: \mathbb{R}^n \rightarrow \mathbb{R}^m$  with  $n$  independent/input variable  $x_i$  and  $m$  dependent/output variable  $y_j$ :
  - Each forward pass of AD is initialized by setting only one of the input variable  $x_i$  and setting the rest to 0 (i.e.,  $\dot{\mathbf{x}} = \mathbf{e}_i$ , where  $\mathbf{e}_i$  is the  $i$ -th unit vector).
  - One execution of forward mode AD computes:  $\dot{y}_j = \frac{\partial y_j}{\partial x_i} |_{\mathbf{x}=\mathbf{a}}, j = 1, \dots, m$
  - Give us one column of the Jacobian matrix at point  $\mathbf{a}$  (the full jacobian can be computed by  $n$  evaluations):

$$J_f = \begin{bmatrix} \frac{\partial y_1}{\partial x_1} & \dots & \frac{\partial y_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial y_m}{\partial x_1} & \dots & \frac{\partial y_m}{\partial x_n} \end{bmatrix} |_{\mathbf{x}=\mathbf{a}}$$

# Reverse Mode AD

- Reverse mode AD propagates derivatives backward from a given output.
- We start by complementing each intermediate variable  $v_i$  with an adjoint (cotangent) representing the sensitivity of a considered output  $y_j$  with respect to changes in  $v_i$ :

$$\bar{v}_i = \frac{\partial y_j}{\partial v_i}$$


- In the first phase, the original function code is run forward, populating intermediate variables  $v_i$  and recording the dependencies in the computational graph.
- In the second phase, derivatives are calculated by propagating adjoints  $\bar{v}_i$  in reverse, from the outputs to the inputs.

Chain rule in the multivariable case:

- $y = f(g_1(x), g_2(x), \dots, g_n(x));$
- $\frac{\partial y}{\partial x} = \sum_{i=1}^n \frac{\partial y}{\partial g_i(x)} \frac{\partial g_i(x)}{\partial x}.$

# Reverse Mode AD:

$$f(x_1, x_2) = \ln(x_1) + x_1 x_2 - \sin(x_2)$$



Forward Primal Trace	
$v_{-1} = x_1$	$= 2$
$v_0 = x_2$	$= 5$
$v_1 = \ln(v_{-1})$	$= \ln(2) = 0.693$
$v_2 = v_{-1} \times v_0$	$= 2 \times 5 = 10$
$v_3 = \sin v_0$	$= \sin 5 = 0.959$
$v_4 = v_1 + v_2$	$= 0.693 + 10$
$v_5 = v_4 - v_3$	$= 10.693 + 0.959$
$y_1 = v_5$	$= 11.652$

The way  $v_{-1}$  influences  $y$  is through  $v_1$  and  $v_2$ :

$$\bar{v}_{-1} = \bar{v}_1 \frac{\partial v_1}{\partial v_{-1}} + \bar{v}_2 \frac{\partial v_2}{\partial v_{-1}}$$

The way  $v_0$  influences  $y$  is through  $v_2$  and  $v_3$ :

$$\bar{v}_0 = \bar{v}_2 \frac{\partial v_2}{\partial v_0} + \bar{v}_3 \frac{\partial v_3}{\partial v_0}$$

$$\bar{v}_4 = \frac{\partial y_1}{\partial v_4} = \frac{\partial y_1}{\partial v_5} \cdot \frac{\partial v_5}{\partial v_4} = \bar{v}_5 \frac{\partial v_5}{\partial v_4}$$


Reverse Adjoint (Derivative) Trace		
$\bar{x}_1 = \bar{v}_{-1}$		$= 5.5$
$\bar{x}_2 = \bar{v}_0$		$= 1.716$
$\bar{v}_{-1} = \bar{v}_{-1} + \bar{v}_1 \frac{\partial v_1}{\partial v_{-1}}$	$= \bar{v}_{-1} + \bar{v}_1 / v_{-1}$	$= 5.5$
$\bar{v}_0 = \bar{v}_0 + \bar{v}_2 \frac{\partial v_2}{\partial v_{-1}}$	$= \bar{v}_0 + \bar{v}_2 \times v_{-1}$	$= 1.716$
$\bar{v}_{-1} = \bar{v}_2 \frac{\partial v_2}{\partial v_{-1}}$	$= \bar{v}_2 \times v_0$	$= 5$
$\bar{v}_0 = \bar{v}_3 \frac{\partial v_3}{\partial v_0}$	$= \bar{v}_3 \times \cos v_0$	$= -0.284$
$\bar{v}_2 = \bar{v}_4 \frac{\partial v_4}{\partial v_2}$	$= \bar{v}_4 \times 1$	$= 1$
$\bar{v}_1 = \bar{v}_4 \frac{\partial v_4}{\partial v_1}$	$= \bar{v}_4 \times 1$	$= 1$
$\bar{v}_3 = \bar{v}_5 \frac{\partial v_5}{\partial v_3}$	$= \bar{v}_5 \times (-1)$	$= -1$
$\bar{v}_4 = \bar{v}_5 \frac{\partial v_5}{\partial v_4}$	$= \bar{v}_5 \times 1$	$= 1$
$\bar{v}_5 = \bar{y}_1$		$= 1$

$$\bar{v}_5 = \bar{y}_1 = \frac{\partial y_1}{\partial y_1} = 1$$

# Reverse Mode AD

- Compute the Jacobian of a function  $f: \mathbb{R}^n \rightarrow \mathbb{R}^m$  with  $n$  independent/input variable  $x_i$  and  $m$  dependent/output variable  $y_j$ .
- An important advantage of the reverse mode is that it is significantly less costly to evaluate (in terms of operation count) than the forward mode for functions with a large number of inputs.
- In the extreme case of  $f: \mathbb{R}^n \rightarrow \mathbb{R}$  only one application of the reverse mode is sufficient to compute the full gradient.
- Because machine learning practice principally involves the gradient of a scalar-valued objective with respect to a large number of parameters, this establishes the reverse mode as the main technique in ML systems.

# Further Reading



- [Automatic differentiation in machine learning: a survey  
\(https://arxiv.org/abs/1502.05767\)](https://arxiv.org/abs/1502.05767)

# Auto-Diff for a Linear Layer



# General Chain Rule

- $y = f(\mathbf{x}): \mathbb{R}^n \rightarrow \mathbb{R};$

- $\nabla f(\mathbf{x}) = \frac{\partial y}{\partial \mathbf{x}} = \left[ \frac{\partial y}{\partial x_1} \quad \dots \quad \frac{\partial y}{\partial x_n} \right] \in \mathbb{R}^n$

- $\mathbf{y} = f(\mathbf{x}): \mathbb{R}^n \rightarrow \mathbb{R}^m;$

- $\frac{\partial \mathbf{y}}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial y_1}{\partial x_1} & \dots & \frac{\partial y_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial y_m}{\partial x_1} & \dots & \frac{\partial y_m}{\partial x_n} \end{bmatrix} \in \mathbb{R}^{m \times n}$

- $\mathbf{y} = f(\mathbf{x}): \mathbb{R}^n \rightarrow \mathbb{R}^m;$
- $\mathbf{z} = g(\mathbf{y}): \mathbb{R}^m \rightarrow \mathbb{R}^k;$
- $\mathbf{z} = f \circ g(\mathbf{x}): \mathbb{R}^n \rightarrow \mathbb{R}^k;$
- $\frac{\partial \mathbf{z}}{\partial \mathbf{x}} = \frac{\partial \mathbf{z}}{\partial \mathbf{y}} \frac{\partial \mathbf{y}}{\partial \mathbf{x}} \in \mathbb{R}^{k \times n}$

# Linear Layer: Forward

- Forward computation of a linear layer:  $\mathbf{Y} = \mathbf{XW}$ 
  - Input:  $\mathbf{X} \in \mathbb{R}^{B \times H_1}$
  - Weight matrix:  $\mathbf{W} \in \mathbb{R}^{H_1 \times H_2}$
  - Output:  $\mathbf{Y} \in \mathbb{R}^{B \times H_2}$
- After the forward pass, we assume that the output will be used in other parts of the model, and will eventually be used to compute a scalar loss  $L \in \mathbb{R}$ .

# Linear Layer: Backward

- During the backward pass through the linear layer, we assume that the derivative  $\frac{\partial L}{\partial \mathbf{Y}} \in \mathbb{R}^{B \times H_2}$  has already been computed and given by:

$$\frac{\partial L}{\partial \mathbf{Y}} = \begin{bmatrix} \frac{\partial L}{\partial Y_{1,1}} & \cdots & \frac{\partial L}{\partial Y_{1,H_2}} \\ \vdots & \ddots & \vdots \\ \frac{\partial L}{\partial Y_{B,1}} & \cdots & \frac{\partial L}{\partial Y_{B,H_2}} \end{bmatrix}$$

- Our goal is to use  $\frac{\partial L}{\partial \mathbf{Y}}$  to compute  $\frac{\partial L}{\partial \mathbf{X}}$  and  $\frac{\partial L}{\partial \mathbf{W}}$ .

# Linear Layer: Backward

- By the general chain rule, we have:

- $\frac{\partial L}{\partial \mathbf{X}} = \frac{\partial L}{\partial \mathbf{Y}} \frac{\partial \mathbf{Y}}{\partial \mathbf{X}}$
- $\frac{\partial L}{\partial \mathbf{W}} = \frac{\partial L}{\partial \mathbf{Y}} \frac{\partial \mathbf{Y}}{\partial \mathbf{W}}$

- But, we do not want to explicitly compute  $\frac{\partial \mathbf{Y}}{\partial \mathbf{X}}$  and  $\frac{\partial \mathbf{Y}}{\partial \mathbf{W}}$ .

- *How can we compute  $\frac{\partial L}{\partial \mathbf{X}}$  and  $\frac{\partial L}{\partial \mathbf{W}}$  without explicitly computing  $\frac{\partial \mathbf{Y}}{\partial \mathbf{X}}$  and  $\frac{\partial \mathbf{Y}}{\partial \mathbf{W}}$ ?*

The Jacobian matrices are too large:  
 $\frac{\partial \mathbf{Y}}{\partial \mathbf{X}} \in \mathbb{R}^{BH_2 \times BH_1}$ ,  $\frac{\partial \mathbf{Y}}{\partial \mathbf{W}} \in \mathbb{R}^{BH_2 \times H_1 H_2}$

# Linear Layer: Backward

- We know that  $\frac{\partial L}{\partial \mathbf{X}}$  should have the same shape as  $\mathbf{X} \in \mathbb{R}^{B \times H_1}$ :

$$\frac{\partial L}{\partial \mathbf{X}} = \begin{bmatrix} \frac{\partial L}{\partial X_{1,1}} & \dots & \frac{\partial L}{\partial X_{1,H_1}} \\ \vdots & \ddots & \vdots \\ \frac{\partial L}{\partial X_{B,1}} & \dots & \frac{\partial L}{\partial X_{B,H_1}} \end{bmatrix}$$

- Let us first try to compute  $\frac{\partial L}{\partial X_{1,1}}$ , by the chain rule, we have:

$$\frac{\partial L}{\partial X_{1,1}} = \sum_{i=1}^B \sum_{j=1}^{H_2} \frac{\partial L}{\partial Y_{i,j}} \frac{\partial Y_{i,j}}{\partial X_{1,1}} = \frac{\partial L}{\partial \mathbf{Y}} \frac{\partial \mathbf{Y}}{\partial X_{1,1}}$$

We have:  $\frac{\partial L}{\partial X_{1,1}} \in \mathbb{R}$ ,  $\frac{\partial L}{\partial \mathbf{Y}} \in \mathbb{R}^{B \times H_2}$ ,  $\frac{\partial \mathbf{Y}}{\partial X_{1,1}} \in \mathbb{R}^{B \times H_2}$ , so this a inner product.

# Linear Layer: Backward

- Since  $\frac{\partial L}{\partial \mathbf{Y}} \in \mathbb{R}^{B \times H_2}$  has already been given, we only need to compute  $\frac{\partial \mathbf{Y}}{\partial X_{1,1}}$

- Recall that  $\mathbf{Y} = \mathbf{XW} = \begin{bmatrix} X_{1,1} & \cdots & X_{1,H_1} \\ \vdots & \ddots & \vdots \\ X_{B,1} & \cdots & X_{B,H_1} \end{bmatrix} \begin{bmatrix} W_{1,1} & \cdots & W_{1,H_2} \\ \vdots & \ddots & \vdots \\ W_{H_1,1} & \cdots & W_{H_1,H_2} \end{bmatrix}$

- $\mathbf{Y} = \begin{bmatrix} \sum_{k=1}^{H_1} X_{1k} W_{k1} & \cdots & \sum_{k=1}^{H_1} X_{1k} W_{kH_2} \\ \vdots & \ddots & \vdots \\ \sum_{k=1}^{H_1} X_{Bk} W_{kH_2} & \cdots & \sum_{k=1}^{H_1} X_{Bk} W_{kH_2} \end{bmatrix}$

- It is easy to check:  $\frac{\partial \mathbf{Y}}{\partial X_{1,1}} = \begin{bmatrix} W_{11} & \cdots & W_{1H_2} \\ 0 & \cdots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \cdots & 0 \end{bmatrix}$

# Linear Layer: Backward

- So the inner product of  $\frac{\partial L}{\partial X_{1,1}} = \frac{\partial L}{\partial Y} \frac{\partial Y}{\partial X_{1,1}}$  can be computed by:

$$\left\langle \begin{bmatrix} \frac{\partial L}{\partial Y_{1,1}} & \cdots & \frac{\partial L}{\partial Y_{1,H_2}} \\ \vdots & \ddots & \vdots \\ \frac{\partial L}{\partial Y_{B,1}} & \cdots & \frac{\partial L}{\partial Y_{B,H_2}} \end{bmatrix}, \begin{bmatrix} W_{1,1} & \cdots & W_{1,H_2} \\ 0 & \cdots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \cdots & 0 \end{bmatrix} \right\rangle = \sum_{k=1}^{H_2} \frac{\partial L}{\partial Y_{1,k}} W_{1k}$$

- Generally, we have  $\frac{\partial L}{\partial X_{i,j}} = \sum_{k=1}^{H_2} \frac{\partial L}{\partial Y_{i,k}} W_{jk}$
- Thus, we have  $\frac{\partial L}{\partial \mathbf{X}} = \frac{\partial L}{\partial \mathbf{Y}} \mathbf{W}^T$

# Linear Layer: Backward

- Using the same strategy of thinking about components one at a time, we can derive a similarly simple equation to compute  $\frac{\partial L}{\partial \mathbf{W}}$  without explicitly forming the Jacobian matrix of  $\frac{\partial \mathbf{Y}}{\partial \mathbf{W}}$ .
- Eventually, we will have  $\frac{\partial L}{\partial \mathbf{W}} = \mathbf{X}^T \frac{\partial L}{\partial \mathbf{Y}}$ .



# Summary of a Linear Layer Computation

- Forward computation of a linear layer:  $\mathbf{Y} = \mathbf{XW}$ 
  - Given input:  $\mathbf{X} \in \mathbb{R}^{B \times D_1}$
  - Given weight matrix:  $\mathbf{W} \in \mathbb{R}^{D_1 \times D_2}$
  - Compute output:  $\mathbf{Y} \in \mathbb{R}^{B \times D_2}$
- Backward computation of a linear layer:
  - Given gradients w.r.t output:  $\frac{\partial L}{\partial \mathbf{Y}} \in \mathbb{R}^{B \times D_2}$
  - Compute gradients w.r.t weight matrix:  $\frac{\partial L}{\partial \mathbf{W}} = \mathbf{X}^T \frac{\partial L}{\partial \mathbf{Y}} \in \mathbb{R}^{B \times D_2}$
  - Compute gradients w.r.t input:  $\frac{\partial L}{\partial \mathbf{X}} = \frac{\partial L}{\partial \mathbf{Y}} \mathbf{W}^T \in \mathbb{R}^{B \times D_2}$

# Linear Layer in PyTorch

```
CLASS torch.nn.Linear(in_features, out_features, bias=True, device=None, dtype=None) \[SOURCE\]
```

Applies a linear transformation to the incoming data:  $y = xA^T + b$ .

This module supports **TensorFloat32**.

On certain ROCm devices, when using float16 inputs this module will use **different precision** for backward.

## Parameters

- **in\_features** (*int*) – size of each input sample
- **out\_features** (*int*) – size of each output sample
- **bias** (*bool*) – If set to `False`, the layer will not learn an additive bias. Default: `True`

## Shape:

- Input:  $(*, H_{in})$  where  $*$  means any number of dimensions including none and  $H_{in} = \text{in\_features}$ .
- Output:  $(*, H_{out})$  where all but the last dimension are the same shape as the input and  $H_{out} = \text{out\_features}$ .

# PyTorch Autograd

# PyTorch Autograd Overview

- PyTorch has a built-in differentiation engine called `torch.autograd`. It supports the automatic computation of gradients for any computational graph.
- Autograd is a reverse automatic differentiation system.
  - Autograd records a graph of the operations (`Function` object);
  - It is a directed acyclic graph (DAG) :
    - Leaves are the input tensors;
    - Roots are the output tensors.
  - By tracing this graph from roots to leaves, automatically compute the gradients using the chain rule.

# PyTorch Autograd Overview

- In a forward pass, **autograd** does two things simultaneously:
  - Run the requested operation to compute a resulting tensor
  - Maintain the operation's gradient function in the DAG.
- The backward pass kicks off when **.backward()** is called on the DAG root. **autograd** then:
  - Compute the gradients from each **.grad\_fn**,
  - Accumulate them in the respective tensor's **.grad** attribute
  - Use the chain rule to propagate all the way to the leaf tensors.

# PyTorch Module

- PyTorch uses modules to represent neural networks. Modules are:
  - Building blocks of stateful computation. PyTorch provides a robust library of modules and makes it simple to define new custom modules, allowing for easy construction of elaborate, multi-layer neural networks.
  - Tightly integrated with PyTorch's autograd system. Modules make it simple to specify learnable parameters for PyTorch's Optimizers to update.
  - Easy to work with and transform. Modules are straightforward to save and restore, transfer between CPU / GPU / TPU devices, prune, quantize, and more.

# PyTorch Module

- Import classes:
  - **Parameter**: A kind of Tensor to be considered a module parameter.
  - **ParameterList**: holds parameters in a list.
  - **ParameterDict**: holds parameters in a dictionary.
  - **Module**: base class for all neural network modules.
  - **Sequential**: a sequential container.
  - **ModuleList**: holds submodules in a list.
  - **ModuleDict**: holds submodules in a dictionary.



# Define A Simple Custom Model

## Define A Simple Custom Model

```
import torch
from torch import nn

class MyLinear(nn.Module):
    def __init__(self, in_features, out_features):
        super().__init__()
        self.weight = nn.Parameter(torch.randn(in_features, out_features))

    def forward(self, input):
        return torch.matmul(input, self.weight)
```

- Inherits from the base Module class. All modules should subclass **Module** for composability with other modules.
- Defines some “state” that is used in computation. Here, the state consists of randomly-initialized weight tensor that define the computation. Because it is defined as a **Parameter**, it is registered for the module and will automatically be tracked and returned from calls to **parameters()**. Parameters can be considered the “learnable” aspects of the module’s computation. Note that modules are not required to have state and can also be stateless.
- Defines a **forward()** function that performs the computation. For this MyLinear module, the input is matrix-multiplied with the weight parameter to produce the output. More generally, the **forward()** implementation for a module can perform arbitrary computation involving any number of inputs and outputs.





# Define A Simple Custom Model

## Define A Simple Custom Model

```
import torch
from torch import nn

class MyLinear(nn.Module):
    def __init__(self, in_features, out_features):
        super().__init__()
        self.weight = nn.Parameter(torch.randn(in_features, out_features))

    def forward(self, input):
        return torch.matmul(input, self.weight)

m = MyLinear(4, 3)
sample_input = torch.randn(4)
print(m(sample_input))

# tensor([-0.7689,  1.0085,  1.7747], grad_fn=<SqueezeBackward4>)
```

- Note that the module itself is *callable*, and that calling it invokes its **forward()** function.
- Each module is associated with “forward pass” and “backward pass”:
  - The “forward pass” applies the computation defined in the module to the given input(s).
  - The “backward pass” computes gradients of module outputs with respect to its inputs. PyTorch’s autograd system automatically takes care of this backward pass computation, so it is *NOT required* to manually implement a **backward()** function for each module.

# Define A Simple Custom Model

- The full set of parameters registered by the module can be iterated through via a call to `parameters()` or `named_parameters()`.

Code	Output
<pre>for parameter in m.named_parameters():     print(parameter)  for parameter in m.parameters():     print(parameter)</pre>	<pre>tensor([-0.3767,  3.0030,  0.0343], grad_fn=&lt;SqueezeBackward4&gt;) (   'weight', Parameter containing:     tensor([[ 0.6101, -0.7031,  1.3140],             [-0.5459, -0.3645,  0.0464],             [ 2.9713,  1.0321,  2.4162],             [ 2.6158, -0.4109,  0.5061]], requires_grad=True))  Parameter containing: tensor([[ -0.6607, -0.3632, -0.7274],         [-0.8555, -1.9544, -0.3640],         [-2.1289,  0.0305,  0.0443],         [-0.6668,  1.7042, -0.0758]], requires_grad=True)</pre>

# Modules as Building Blocks

## Modules as Building Blocks

```
net = nn.Sequential(  
    MyLinear(4, 3),  
    nn.ReLU(),  
    MyLinear(3, 1)  
)  
  
sample_input = torch.randn(4)  
print(net(sample_input))  
  
# tensor([-0.6749], grad_fn=<AddBackward0>)
```

- Modules can contain other modules, making them useful building blocks for developing more elaborate functionality. The simplest way to do this is using the **nn.Sequential** module. It allows us to chain together multiple modules.
- **nn.Sequential** automatically feeds the output of the first **MyLinear** module as input into the **ReLU**, and the output of that as input into the second **MyLinear** module. As shown, it is limited to in-order chaining of modules with a single input and output.

# Modules as Building Blocks

## Modules as Building Blocks

```
class Net2(nn.Module):
    def __init__(self):
        super().__init__()
        self.layer0 = MyLinear(4, 3)
        self.layer1 = MyLinear(3, 1)

    def forward(self, x):
        x = self.layer0(x)
        x = F.relu(x)
        x = self.layer1(x)
        return x

net2 = Net2()
sample_input = torch.randn(4)
print(net2(sample_input))

# tensor([-0.2827], grad_fn=<SqueezeBackward4>)
```

- In general, it is recommended to define a custom module for anything beyond the simplest use cases, as this gives full flexibility on how submodules are used for a module's computation.
- This module is composed of two “children” or “submodules” (**layer0** and **layer1**) that define the layers of the neural network and are utilized for computation within the module's **forward( )** method.

# Modules as Building Blocks

`children()` [\[SOURCE\]](#)

Return an iterator over immediate children modules.

**Yields**

*Module* – a child module

**Return type**

*Iterator*[*Module*]

`modules()` [\[SOURCE\]](#)

Return an iterator over all modules in the network.

**Yields**

*Module* – a module in the network

**Return type**

*Iterator*[*Module*]

## Code

```
print("Check net.children")
for child in net.children():
    print(child)
```

```
print("Check net.modules")
for child in net.modules():
    print(child)
```

## Output

```
<--Check net.children()-->
MyLinear()
ReLU()
MyLinear()

<--Check net.modules()-->
Sequential(
  (0): MyLinear()
  (1): ReLU()
  (2): MyLinear()
)
MyLinear()
ReLU()
MyLinear()
```

# Modules as Building Blocks

`children()` [\[SOURCE\]](#)

Return an iterator over immediate children modules.

**Yields**

*Module* – a child module

**Return type**

*Iterator*[*Module*]

`modules()` [\[SOURCE\]](#)

Return an iterator over all modules in the network.

**Yields**

*Module* – a module in the network

**Return type**

*Iterator*[*Module*]

## Code

```
print("<--Check net2.children()-->")
for child in net2.children():
    print(child)

print("<--Check net.modules()-->")
for child in net2.modules():
    print(child)
```

## Output

```
<--Check net2.children()-->
MyLinear()
MyLinear()

<--Check net2.modules()-->
Net2(
  (layer0): MyLinear()
  (layer1): MyLinear()
)
MyLinear()
MyLinear()
```

# Modules as Building Blocks

`named_children()` [\[SOURCE\]](#)

Return an iterator over immediate children modules, yielding both the name of the module as well as the module itself.

**Yields**

`(str, Module)` – Tuple containing a name and child module

**Return type**

`Iterator[Tuple[str, Module]]`

`named_modules(memo=None, prefix='', remove_duplicate=True)` [\[SOURCE\]](#)

Return an iterator over all modules in the network, yielding both the name of the module as well as the module itself.

**Parameters**

- **memo** (*Optional*[*Set*[*Module*]]) – a memo to store the set of modules already added to the result
- **prefix** (*str*) – a prefix that will be added to the name of the module
- **remove\_duplicate** (*bool*) – whether to remove the duplicated module instances in the result or not

**Yields**

`(str, Module)` – Tuple of name and module

## Code

```
print("<--Check net2.named_children()-->")
for child in net2.named_children():
    print(child)

print("<--Check net.named_modules()-->")
for child in net.named_modules():
    print(child)
```

## Output

```
<--Check net2.named_children()-->
('layer0', MyLinear())
('layer1', MyLinear())

<--Check net.named_modules()-->
('', Net2(
  (layer0): MyLinear()
  (layer1): MyLinear()
))
('layer0', MyLinear())
('layer1', MyLinear())
```

# How Does Autograd Execute?

## Code

```
x = torch.randn(4)
x.requires_grad_(True)
x.retain_grad()
z = net2(x)
z.retain_grad()

print("x:", x)
print("w0:", net2.layer0.weight)
print("w1:", net2.layer1.weight)
print("z:", z)

z.backward()

print("dz:", z.grad)
print("dw1:",
      net2.layer1.weight.grad)
print("dw2:",
      net2.layer0.weight.grad)
print("dx:", x.grad)
```

## Output

```
x: tensor([ 0.9785,  0.4565, -0.4396, -0.1090],
          requires_grad=True)
w0: Parameter containing:
tensor([[ -0.4769, -0.9407,  1.1517],
        [ -0.4809,  1.2256,  0.9053],
        [  1.5409, -1.2598, -3.7088],
        [ -1.0534, -0.5934,  0.0647]]),
       requires_grad=True)
w1: Parameter containing:
tensor([[ 1.7888],
        [ 0.1920],
        [-1.8594]]), requires_grad=True)
z: tensor([-5.8328], grad_fn=<SqueezeBackward4>)
dz: tensor([1.])
dw1: tensor([[0.0000],
             [0.2575],
             [3.1636]])
dw2: tensor([[ 0.0000,  0.1879, -1.8195],
             [ 0.0000,  0.0877, -0.8488],
             [-0.0000, -0.0844,  0.8174],
             [-0.0000, -0.0209,  0.2027]])
dx: tensor([-2.3221, -1.4479,  6.6540, -0.2343])
```

### TORCH.TENSOR.BACKWARD

`Tensor.backward(gradient=None, retain_graph=None, create_graph=False, inputs=None)`[\[source\]](#)

Computes the gradient of current tensor wrt graph leaves.

The graph is differentiated using the chain rule. If the tensor is non-scalar (i.e. its data has more than one element) and requires gradient, the function additionally requires specifying `gradient`. It should be a tensor of matching type and location, that contains the gradient of the differentiated function w.r.t. `self`.

This function accumulates gradients in the leaves - you might need to zero `.grad` attributes or set them to `None` before calling it. See [Default gradient layouts](#) for details on the memory layout of accumulated gradients.


### TORCH.TENSOR.RETAIN\_GRAD

`Tensor.retain_grad()` → [None](#)

Enables this Tensor to have their `grad` populated during `backward()`. This is a no-op for leaf tensors.



# How Does Autograd Execute?



# Locally Disabling Gradient Computation

- The most fine-grained exclusion of subgraphs from gradient computation is setting the `requires_grad` field of a tensor.
- To disable gradients across entire blocks of code, there are also context managers:
  - no-grad mode
  - inference mode.

# Setting `requires_grad`

- Recall that “A **`torch.Tensor`** is a multi-dimensional matrix containing elements of a single data type”.
- The attribute **`requires_grad`** is a flag, that allows for fine-grained exclusion of subgraphs from gradient computation.
- The attribute **`requires_grad`** is set to false by default unless wrapped in a **`nn.Parameter`**.
- It takes effect in both the forward and backward passes:
  - During the forward pass, an operation is only recorded in the backward graph if at least one of its input tensors require grad.
  - During the backward pass (**`.backward()`**), only leaf tensors with **`requires_grad=True`** will have gradients accumulated into their **`.grad`** fields.

# Setting `requires_grad`

- Although every tensor has this flag, setting it only makes sense for *leaf tensors* (tensors that do not have a `grad_fn`, e.g., a `nn.Module`'s `parameters`).
- Non-leaf tensors (tensors that do have `grad_fn`) are tensors that have a backward graph associated with them, so their gradients will be needed as an intermediary result to compute the gradient for a leaf tensor that requires grad. Thus, all non-leaf tensors will automatically have `requires_grad=True`.
- Setting `requires_grad` should be the main way you control which parts of the model are part of the gradient computation.
  - E.g, if you need to freeze parts of your pretrained model during model fine-tuning.
  - To freeze parts of your model, simply apply `.requires_grad_(False)` to the parameters that you don't want updated.
  - Since computations that use these parameters as inputs would not be recorded in the forward pass, they won't have their `.grad` fields updated in the backward pass because they won't be part of the backward graph in the first place, as desired.
- `requires_grad` can also be set at the module level with `nn.Module.requires_grad_()`. When applied to a module, `.requires_grad_()` takes effect on all of the module's parameters

# Setting `requires_grad`

```
requires_grad_(requires_grad=True) \[SOURCE\]
```

Change if autograd should record operations on parameters in this module.

This method sets the parameters' `requires_grad` attributes in-place.

This method is helpful for freezing part of the module for finetuning or training parts of a model individually (e.g., GAN training).

See [Locally disabling gradient computation](#) for a comparison between `.requires_grad_()` and several similar mechanisms that may be confused with it.

## Parameters

**`requires_grad`** (*bool*) – whether autograd should record operations on parameters in this module.  
Default: `True`.

## Returns

`self`

## Return type

`Module`

# Setting requires\_grad



## Code

```
net2.layer1.requires_grad_(False)

x = torch.randn(4)
x.requires_grad_(True)
x.retain_grad()
z = net2(x)
z.retain_grad()

print("x:", x)
print("w0:", net2.layer0.weight)
print("w1:", net2.layer1.weight)
print("z:", z)

z.backward()

print("dz:", z.grad)
print("dw1:",
      net2.layer1.weight.grad)
print("dw2:",
      net2.layer0.weight.grad)
print("dx:", x.grad)
```

## Output

```
x: tensor([-0.3019,  0.3072, -0.9097, -2.0537],
           requires_grad=True)
w0: Parameter containing:
tensor([[ -1.9223, -0.5765, -0.3053],
        [  1.1695, -0.1050,  1.2679],
        [  1.0916,  0.5429, -0.6907],
        [-1.3432,  0.0669,  1.4241]],
        requires_grad=True)
w1: Parameter containing:
tensor([[0.5816],
        [0.5796],
        [2.0542]])
z: tensor([1.5735], grad_fn=<SqueezeBackward4>)
dz: tensor([1.])
dw1: None
dw2: tensor([[ -0.1756, -0.0000, -0.0000],
             [  0.1787,  0.0000,  0.0000],
             [-0.5291, -0.0000, -0.0000],
             [-1.1945, -0.0000, -0.0000]])
dx: tensor([-1.1181,  0.6803,  0.6349, -0.7813])
```

## TORCH.TENSOR.BACKWARD

`Tensor.backward(gradient=None, retain_graph=None, create_graph=False, inputs=None)`[\[source\]](#)

Computes the gradient of current tensor wrt graph leaves.

The graph is differentiated using the chain rule. If the tensor is non-scalar (i.e. its data has more than one element) and requires gradient, the function additionally requires specifying `gradient`. It should be a tensor of matching type and location, that contains the gradient of the differentiated function w.r.t. `self`.


This function accumulates gradients in the leaves - you might need to zero `.grad` attributes or set them to `None` before calling it. See [Default gradient layouts](#) for details on the memory layout of accumulated gradients.

## TORCH.TENSOR.RETAIN\_GRAD

`Tensor.retain_grad()` → `None`

Enables this Tensor to have their `grad` populated during `backward()`. This is a no-op for leaf tensors.

# Setting requires\_grad



# Grad Modes

Mode	Excludes operations from being recorded in backward graph	Skips additional autograd tracking overhead	Tensors created while the mode is enabled can be used in grad-mode later	Examples
default			✓	Forward pass
no-grad	✓		✓	Optimizer updates
inference	✓	✓		Data processing, model evaluation



# Default Mode

- The “default mode” is the mode we are implicitly in when no other modes like no-grad and inference mode are enabled.
- The most important thing to know about the default mode is that it is the only mode in which `requires_grad` takes effect. `requires_grad` is always overridden to be `False` in the two other modes.



# No-grad Mode

- Computations in no-grad mode behave as if none of the inputs require grad.
- Computations in no-grad mode are never recorded in the backward graph even if there are inputs that have `require_grad=True`.
- Enable no-grad mode when you need to perform operations that should not be recorded by `autograd`, but you'd still like to use the outputs of these computations in grad mode later.
- This context manager makes it convenient to disable gradients for a block of code or function without having to temporarily set tensors to have `requires_grad=False`, and then back to `True`.

# Recall Our First Training Script

## Define the test function

```
def test_loop(dataloader, model, loss_fn):
    model.eval()
    size = len(dataloader.dataset)
    num_batches = len(dataloader)
    test_loss, correct = 0, 0

    with torch.no_grad():
        for X, y in dataloader:
            pred = model(X)
            test_loss += loss_fn(pred, y).item()
            correct += (pred.argmax(1) == y).type(torch.float).sum().item()
    test_loss /= num_batches
    correct /= size
    print(f"Test Error: \n Accuracy: {(100*correct):>0.1f}%, Avg loss: {test_loss:>8f} \n")
```

# Inference Mode

- Inference mode is the extreme version of no-grad mode.
- Computations in inference mode are not recorded in the backward graph, but enabling inference mode will allow PyTorch to speed up your model even more.
- Drawback: tensors created in inference mode will not be able to be used in computations to be recorded by **autograd** after exiting inference mode.
- Enable inference mode when you are performing computations that don't need to be recorded in the backward graph, **AND** you don't plan on using the tensors created in inference mode in any computation that is to be recorded by **autograd** later.

# Compare Different Modes



## Compare different modes

```
dim1 = 4096
dim2 = 8192

class MyLinear(nn.Module):
    def __init__(self, in_features, out_features):
        super().__init__()
        self.weight = nn.Parameter(torch.randn(in_features, out_features))

    def forward(self, input):
        return torch.matmul(input, self.weight)

class Net3(nn.Module):
    def __init__(self):
        super().__init__()
        self.layer0 = MyLinear(dim1, dim2)
        self.layer1 = MyLinear(dim2, dim2)
        self.layer2 = MyLinear(dim2, dim2)
        self.layer3 = MyLinear(dim2, 1)

    def forward(self, x):
        x = self.layer0(x)
        x = F.relu(x)
        x = self.layer1(x)
        x = F.relu(x)
        x = self.layer2(x)
        x = F.relu(x)
        x = self.layer3(x)
        return x

net3 = Net3()
x = torch.randn(256, dim1)
```

Code	Output
<pre>start_time = time.time() z = net3(x) end_time = time.time() print("Forward computation takes:", end_time-start_time)</pre>	Forward computation takes: 0.5174 second
<pre>start_time = time.time() with torch.no_grad():     z = net3(x) end_time = time.time() print("Forward computation takes: ", end_time-start_time)</pre>	Forward computation takes: 0.4768 second
<pre>start_time = time.time() with torch.inference_mode():     z = net3(x) end_time = time.time() print("Forward computation takes: ", end_time-start_time)</pre>	Forward computation takes: 0.4341 second

# References

- [Automatic differentiation in machine learning: a survey \(https://arxiv.org/abs/1502.05767\)](https://arxiv.org/abs/1502.05767)
- <http://cs231n.stanford.edu/handouts/linear-backprop.pdf>
- <https://pytorch.org/docs/stable/generated/torch.nn.Linear.html>
- <https://pytorch.org/docs/stable/notes/autograd.html>
- <https://pytorch.org/docs/stable/notes/modules.html>